

Assignment 0

git'n Started

Prof. Darrell D. E. Long
CSE 13S – Winter 2022

Due: January 6th at 11:59 pm

1 Introduction

Dire Straits is a great band. Someone tells you they like 'Brothers in Arms' and immediately you know they're a stupid annoying git.

—Alexei Sayle

The aim of this first assignment will be for you to set up your GitLab repositories and gain an understanding of how `git` works. We will review several `git` commands that you will help you in the long run. This document will be helpful for troubleshooting `git` issues in the future and also includes the submission policy. You will find this quite helpful in the future if you ever have any issues with `git` or submitting. *Ideally*, this assignment should be completed during your discussion section.

2 Source Code Control

The deliverables for each of your assignments will be maintained through your `git` repository. We are using GitLab, which is a service coupled with the version-control capabilities of `git`. `git` allows you to maintain multiple versions of your source files, also known as version control. Version control is the practice of tracking and making changes to code, such that in the event of some accident while coding, it is always possible to restore your code to a previous state. `git` is used through a set of commands within a repository, a version-controlled directory that stores your files.

2.1 Setting Up SSH Keys

You will first need to *clone* your GitLab repository. It is highly recommended that you use `git` over SSH rather than HTTP. SSH, the *Secure Socket Shell*, is a protocol that provides secure network communication. Using `git` over SSH allows for user authentication using SSH keys, ridding the need to enter your username or password each time you want to push or pull changes to your GitLab repository. Other tools that use SSH include `scp` and `sftp`. The former is a *Secure Copy Protocol* designed to transfer files to and from servers over an SSH connection. The latter is a *SSH File Transfer Protocol* and is typically used as a remote file system protocol.

SSH keys come in *pairs*: a *private* key and a *public* key. Data encrypted with some public key can only be decrypted with the corresponding private key and vice versa; public key cryptography. You may

find it useful to keep a copy of your SSH keys somewhere — like a USB stick. As long as you have access to the keys, you can securely authenticate with git. To generate an SSH keypair:

```
$ ssh-keygen
```

This assumes the use of UNIX, or on macOS — attend section if you wish to learn how to generate a key pair on Windows.

The SSH keypair generated by the default prompt answers will be sufficient for your needs for use with GitLab. These keys last a long time, so try not to lose them. Make sure you add the *public* key of the generated key pair to GitLab and that it is an RSA key. To print your public key so that you can copy it to your clipboard, enter the following:

```
$ cat ~/.ssh/id_rsa.pub
```

This specific command uses the UNIX utility cat, which is designed to concatenate and print files. The argument supplied to cat is the path to your public key. All keys generated by ssh-keygen reside under the ssh directory. After adding the key, you will be ready to clone your GitLab repository. For more in-depth instructions on generating and adding SSH keys, as well as other GitLab basics, please refer to this link:

<https://git.ucsc.edu/help/gitlab-basics/README.md>

2.2 Cloning Your Repository

To clone your repository, run the following command, substituting <CruzID> with your CruzID:

```
$ git clone git@git.ucsc.edu:cse13s/winter2022/<CruzID>.git cse13s
```

You will be prompted for permission to authenticate with the server. When permitted, the command will clone your repository onto your machine into a directory named cse13s in the current working directory. Use the cd command to enter the asgn0 directory in your cloned cse13s repository to start your work for assignment 0.

```
$ cd cse13s/asgn0
```

2.3 The .ssh Directory and File Permissions

We'll take a moment here to discuss the contents of the ~/.ssh directory found on macOS and UNIX-based systems ¹, along with basic file permissions on UNIX. There are two files of note, other than generated public and private keys: ~/.ssh/known_hosts and ~/.ssh/authorized_keys.

The ~/.ssh/known_hosts file contains the SSH fingerprints of every remote machine you choose to SSH onto. Whenever you try to SSH onto a new machine, a prompt will appear asking if you want to continue connecting, where an affirmative response will cause the remote machine's SSH fingerprint to be *appended* to the known hosts file.

¹The ~/.ssh directory equivalent on Windows is typically C:\Users\<username>\.ssh. Your coursework, however, should be done on Ubuntu 20.04 or later.

The `~/.ssh/authorized_keys` file contains the public keys of all remote clients that are authorized to remotely log onto the server using SSH authentication, the server in which is the machine that contains the file of authorized keys.

Each of the aforementioned SSH-related files, including the public and private keys, require specific *file permissions*. File permissions on UNIX are what allow you as a user to read and modify (write) files, as well as what prevents others from being able to read and modify files. File permissions are split into three access groups:

1. User — the owner of the file, typically the user that created the file.
2. Group — a specific set of users.
3. Other — anyone who isn't the user or part of the authorized group.

These three access groups each have three access modes:

1. Read — the ability to view the contents of a file.
2. Write — the ability to modify or delete the contents of a file.
3. Execute — the ability to run a file as a program.

Note that the three access modes were described with files in mind. UNIX directories are also considered files, but the access modes do slightly different things. Read permission for a directory allows permitted users to list files in the directory with a program such as `ls`. Write permission for a directory allows permitted users to modify, create, or delete files in the directory. Lastly, execute permission for a directory allows permitted users to access files within the directory, as well as make the directory their current working directory by using a shell builtin such as `cd` or `pushd`.

The granting and removal of file permissions can be done using the `chmod` utility, typically specifying the desired permissions with three *octal* digits. Why octal? There are exactly two states that each permission can be in: read permission is granted, or it isn't. As a bit, this means either a 0 or 1. Since there are read, write, and execute permissions, this naturally means $2^3 = 8$ possible states, which can be perfectly expressed using one octal digit. Since there are three different access groups (user, group, and other), three octal digits are used. Consider the octal value `68`. In binary, this is `1102`. The read bit is the most significant bit, the execute bit is the least significant bit, and the write bit is between the two. For this specific example, only the read and write bits are set.

The permissions of your `~/.ssh` directory should be `700`. The leftmost octal digit signifies user permissions, the middle octal digit signifies group permissions, and the rightmost octal digit signifies other permissions. Thus, the permissions on this directory allows only the user to enter it, as well as read and modify files. To explicitly set the permissions as `700` using `chmod`:

```
$ chmod 700 ~/.ssh
```

Read the man page for `chmod`, then make sure the permissions for the files in your SSH directory are set as follows:

- `known_hosts` — `644`
- `authorized_keys` — `600`

- any private key — 600
- any public key — 644

You can create as many public and private keys with `ssh-keygen` as you want, but you really *only need one*.

3 Hello World!

You will be creating a simple C program which will simply print “Hello World!” **You can find also find a tutorial of this program in Chapter 1 §1.1** in your textbook, *The C Programming Language* by Kernighan & Ritchie.

1. Make sure you are in the correct directory: `asgn0`. You can check your *current working directory* using this command:

```
$ pwd
```

2. Create the program source `hello.c` with your text editor of choice. This means text editors such as `vi` and `emacs`. Notepad and Word are *not* text editors. To open up `hello.c` for editing with `vi`:

```
$ vi hello.c
```

It should be noted that the `vim` text editor has largely succeeded `vi`. In fact, many systems simply have `vi` aliased to `vim`. You can check if your version of `vi` is an alias by running:

```
$ vi --version
```

Using either `vi` or `vim` is perfectly acceptable for your assignments, but you may find `vim` easier to get accustomed to with all the quality-of-life improvements that has been over the years.

3. Include the header for the `<stdio.h>` library. This is needed by the `printf()` function that prints formatted strings to `stdout`, what you think of as the console.

```
hello.c
1 #include <stdio.h>
```

4. Type your `main()` function. Every C program *must* have a `main()` function which returns an `int`. A return value of 0 indicates program success, and a non-zero return indicates the occurrence of some error.

```
hello.c
1 #include <stdio.h>
2
3 int main(void) {
4     return 0;
5 }
```

5. In `main()` (between the curly braces) is where you will type the print statement. It is *crucial* that your print statement matches the one given here. **You will be docked points otherwise.**

```
hello.c
1 #include <stdio.h>
2
3 int main(void) {
4     printf("Hello World!\n");
5     return 0;
6 }
```

6. Save your work and exit your text editor to return to the command line. With `vi` this means entering normal mode by hitting `esc` and entering the command (indicated with a prefixed colon) “`:wq`” to save and quit. `ZZ` works as well. You are allowed to use `vi`, `vim`, `nvim`, `emacs`, `atom`, `Sublime` — but *not* a Windows editor. *Editing* the file using Windows — or `Notepad`, or even `Word` — is grounds for receiving a zero.
7. You should now be back on the command line. You should now compile and run your code to verify its correctness. To compile your code, run:

```
$ clang -Wall -Wextra -Werror -Wpedantic -o hello hello.c
```

This will compile your code with the compiler flags required by the class. `clang` is the C compiler that we will be using — not `gcc`, not `cc`. You *must* use `clang`. The `-Wall -Wextra -Werror -Wpedantic` arguments are the set of compiler flags you must use when compiling your code. This specific set of compiler flags is commonly referred to as the “take no prisoners” compiler flags. Simply put, together they catch pretty much everything that a compiler can catch (there are a few more esoteric warnings that can be enabled). Here are some links for you to investigate what each flag does:

<https://releases.llvm.org/10.0.0/tools/clang/docs/UsersManual.html>

<https://releases.llvm.org/10.0.0/tools/clang/docs/DiagnosticsReference.html>

The provided links are for version 10 of `clang`, but it is fine if you have version 13.

8. If you’ve done everything correctly up to this point, the compilation process should run silently and return no errors. However, if you do run into any errors, lab sections, and Piazza will be your best friends. Resist the urge to immediately use Google.
9. After successfully compiling your program, there should now be an executable file named `hello` in the current working directory. To list out all the files in the current working directory use `ls`:

```
$ ls
```

10. To display the current working directory, or what you think of as the directory you’re currently in, use `pwd`:

```
$ pwd
```

11. To run the `hello` program, enter:

```
$ ./hello
```

12. The `.` (usually called “dot”) refers to the *current working directory*. Your shell has a `PATH` environment variable, a colon-delimited list of directories that it looks through when you enter a command. Since your current working directory should never be in your `PATH`, you must specify the directory that your program can be found in order to run it. If the output of running your program is correct, you should then submit your working *source code* to git. You should submit source code *only*: no executables. What is a `PATH`? It is an environment variable whose value is a colon-delimited list of directory names. Each directory in this list indicates a location where executable programs can be found. Directories that are usually included in the `PATH` include `/usr` and `/usr/bin`. The current working directory should *never* be added to the `PATH` since doing so would be a serious vulnerability. Imagine some adversary managed to sneak in a compromised `ls` binary into your current working directory. If your current working directory was first in the `PATH`, then the compromised `ls` would be found and run first, most likely leaving your machine in an undesirable state.

```
$ git add hello.c
$ git commit -m "Adding finished hello.c"
$ git push
```

The above three commands will add, commit, and push `hello.c` to git. In-depth description of each of these commands will be provided in the following section. To verify that `hello.c` was added, check your repository:

<https://git.ucsc.edu/gitlab/cse13s/winter2022/<CruzID>/asgn0>

Only in this case do you perform one commit at the end. In general, you should commit after every significant change. **Warning: you should *never* push binary files. This includes executable programs and object files, as well as any files generated during the compilation process of a program.**

13. The only other file to be submitted for assignment 0 is your signed `CHEATING.pdf`. This file can be found on Canvas and/or under Piazza resources. **Note: You do not create a PDF file by simply appending `.pdf` to its name.** You will submit `CHEATING.pdf` the same way you did `hello.c`: adding, committing, then pushing.

4 Useful Git Commands

The following commands are used through `git` for version control. For this assignment, you will have used the `clone`, `add`, and `push` commands. This section will serve as a brief description and use of frequently used `git` commands that you will most likely use throughout the quarter, if not your entire career as a computing professional.

4.1 git config

This command lets you set configuration variables that tune git to operate the way you want it to. The main things you will likely want to do when you get started with git are establishing your identity, as well as the default text editor when typing up longer commits.

Perform the following command to set your name and email address. Notice the `--global` in the command. That is a *command-line option* and indicates that the following configuration should be used globally in every git repository you have. Unless otherwise specified, configurations by default are applied only to the local repository.

```
$ git config --global user.name "<your name>"
$ git config --global user.email "<your email>"
```

To set your default editor as vim:

```
$ git config --global core.editor vim
```

To check all the configurations simply run:

```
$ git config --list
```

To check the value of a specific key, or setting, just supply it as the sole argument after `git config`. For instance, to check the configured email:

```
$ git config user.email
```

4.2 git help

When starting out with git, you may find yourself frequently needing to refresh your memory on certain commands. The command `git help` will prove invaluable in this regard. There are three ways to display the man page for any git command:

```
$ git help <command>
$ git <command> --help
$ man git-<command>
```

For example, to view the man page for `git clone`, the subject of the next section, any of the following can be run:

```
$ git help clone
$ git clone --help
$ man git-clone
```

A man page (short for manual page) is software documentation for tools and programs found on UNIX systems. To view a man page:

```
$ man <function, program, tool>
```

These manual pages are typically divided into sections, depending on their respective purposes. General commands are found in section 1, system calls in section 2, and library functions, such as the `printf()` function used in this assignment, are found in section 3. So, to view the man page for `printf()`:

```
$ man 3 printf
```

4.3 git clone

This command clones a repository from a server onto your local machine. This downloads a copy of the repository which is stored on a server for local editing. Meaning, any changes that need to be sent back to the server will need to be *added*, *committed* and *pushed*. Here is an example of cloning over ssh:

```
$ git clone user@somemachine:path/to/repo
```

4.4 git add

This command allows you to add files into your repository and stages them to the git source tree. Any file that has been changed since the time it was last added needs to be added again.

```
$ git add file1 file2
```

Keep in mind, adding files with this command does *not* commit them. You still need to commit the changes with the `git commit` command.

4.5 git commit

This command creates a checkpoint for each file which was added using the previous command, `git add`. You can think of it like capturing a snapshot of the current staged changes. These snapshots are then safely committed. Each commit has a unique commit ID along with a message about the commit.

```
$ git commit -m "A short informative message about any changes"
```

To commit all the changed files, you can use the command `git commit -a` which can also be combined with the `-m` option. This will only commit files that have been added and committed at least once before. Without the `-m` flag, you will be taken into the default git editor to enter your commit message. A forewarning: don't commit rude comments — the TAs and graders will see them.

You should commit working versions of your code frequently so in the case you mess something up, like accidentally deleting your code, you can use `git checkout HEAD` to revert to the most recent commit.

4.6 git checkout

This command allows you to set the state of your repository to the state of your repository at the time of a different commit. The reverting of state can be performed per file, meaning that you can use `git checkout` to restore a specific file to its state in a different commit. To checkout a commit:

```
$ git checkout <commit>
```

To checkout restore a file to its state at a different commit:

```
$ git checkout <commit> -- <file>
```

This last command also works to retrieve files that you may have accidentally deleted locally. This alone should provide good incentive to add, commit, and push changes to your files often.

4.7 git log

This command provides a list of the commits that have been made on the repository. It provides access to look up commit times, messages, and IDs.

```
$ git log
```

4.8 git push

This command pushes all of your local commits to the upstream repository. It pushes all of your changes to the directory which is stored on-line. You *must* do this to turn in your work for this class. If you do not run this command after committing, *none* of your work will be turned in.

```
$ git push
```

4.9 git pull

This command fetches and downloads content from a remote repository. Your local repository is immediately updated to match the fetched content. `git pull` is actually a combination of `git fetch` followed by `git merge`. The first half of `git pull` will execute `git fetch` on the local branch that HEAD is pointed at. After the contents are fetched, the second half of `git pull` will merge the work-flow creating a new merge commit ID and HEAD is updated to point to the new commit.

```
$ git pull
```

4.10 git ls-files

This command lists all files in the current directory that have been checked into the repository. This will be useful for making sure you have submitted all required deliverables for each assignment.

```
$ git ls-files
```

4.11 git status

This command provides a status of which files have been added and staged for the next commit, as well as unpushed changes.

```
$ git status
```

5 Honesty

Academic honesty is very important in computer science, and life in general. The goal of this course is for you to learn the material, not simply for you to get a mark on your transcript saying you passed the class. All students in the class must sign and turn in an acknowledgment that they understand the cheating policy for the class. **We will not accept or grade any assignments from a student unless they have turned in the CHEATING.pdf.** We encourage you to ask for clarifications in the academic policy if you have any questions.

6 Deliverables

For this class, you will be turning in all of your work through git. All the files you need to turn in for an assignment will be found and listed in the Deliverables section of every assignment PDF. Files will need to be added to the corresponding assignment directory, committed, and pushed.

You will need to turn in:

1. CHEATING.pdf
2. hello.c

7 Submission

The cost of freedom is always high, but Americans have always paid it. And one path we shall never choose, and that is the path of surrender, or submission.

—John F. Kennedy

Now that you have learned about some useful git commands, it's time to put them to use. The steps to submitting assignments will not change throughout the course. If you ever forget the steps, refer back to this PDF. Remember: *add*, *commit*, and *push*! In the case you do mess something up, *don't panic*. Take a step back and think things throughly. The Internet, TAs and tutors are here as resources.

1. Add it!

```
$ git add CHEATING.pdf hello.c
```

As mentioned before, you will need to first add the files to your repository using the `git add <filenames>` command. You will be submitting these files into the `asgn0` directory.

2. Commit it!

```
$ git commit -m "Your commit message here"
```

Changes to these files will be committed to the repository with `git commit`. The command should also include a commit message describing what changes are included in the commit.

3. Push it!

```
$ git push
```

The committed changes are then sync'd up with the remote server using the `git push` command. You must be sure to push your changes to the remote server or else they will not be received by the graders.

4. Submit the commit ID on Canvas! You can find the most recent commit ID by using `git log` or searching for it through the GitLab web interface. **Your assignment is turned in *only* after you have pushed and submitted the commit ID you want graded on Canvas. “I forgot to push” and “I forgot to submit my commit ID” are not valid excuses. It is *highly* recommended to commit and push your changes *often*.**

8 Supplemental Readings

- *Version Control with Git* by Loeliger & McCullough ← Read this! Now!
 - Chapter 3 – Getting Started (pg. 22–25)
- *The C Programming Language* by Kernighan & Ritchie ← It is a *huge* mistake to not read this!
 - Chapter 1 §1.1
- *vi and Vim Editors* by Robbins & Lamb
 - Chapter 1 §1.4 & §1.5



Programming in C is like giving a monkey a chainsaw.